

Introduction to Julia

Reference: [The Julia Documentation](#), [The Julia–Matlab–Python Cheatsheet](#)

This lecture gives an overview of Julia. In these notes we focus on the aspects of Julia and computing that are essential to numerical computing:

1. Integers: We discuss briefly how to create and manipulate integers, and how to see the underlying bit representation.
2. Strings and parsing: We discuss how to create and manipulate strings and characters, and how we can convert a string of 0's and 1's to an integer or other type.
3. Vectors and matrices: We discuss how to build and manipulate vectors and matrices (which are both types of *arrays*). Later lectures will discuss linear algebra.
4. Types: In Julia everything has a type, which plays a similar role to classes in Python. Here we discuss how to make new types, for example, a complex number in radial format.
5. Loops and branches: We discuss `if`, `for` and `while`, which work similar to Python.
6. Functions: We discuss the construction of named and anonymous functions. Julia allows overloading functions for different types, for example, we can overload `*` for our radial complex type.
7. Modules, Packages, and Plotting: We discuss how to load external packages, in particular, for plotting.

Note some subsections are labeled *advanced*, these are non-examinable and will not be necessary in problem sheets or exams.

1. Integers

Julia uses a math-like syntax for manipulating integers:

```
In [2]: 1 + 1 # Addition
```

```
Out[2]: 2
```

```
In [2]: 2 * 3 # Multiplication
```

```
Out[2]: 6
```

```
In [3]: 2 / 3 # Division
```

```
Out[3]: 0.6666666666666666
```

```
In [4]: x = 5; # semicolon is optional but supresses output if used in the last line
        x^2 # Powers
```

```
Out[4]: 25
```

In Julia everything has a type. This is similar in spirit to a class in Python, but much more

lightweight. An integer defaults to a type `Int`, which is either 32-bit (`Int32`) or 64-bit (`Int64`) depending on the processor of the machine. There are also 8-bit (`Int8`), 16-bit (`Int16`), and 128-bit (`Int128`) integer types, which we can construct by converting an `Int`, e.g. `Int8(3)`.

These are all "primitive types", instances of the type are stored in memory as a fixed length sequence of bits. We can find the type of a variable as follows:

```
In [5]: typeof(x)
```

```
Out[5]: Int64
```

For a primitive type we can see the bits using the function `bitstring`:

```
In [6]: bitstring(Int8(1))
```

```
Out[6]: "00000001"
```

Negative numbers may be surprising:

```
In [7]: bitstring(-Int8(1))
```

```
Out[7]: "11111111"
```

This is explained in detail in Chapter [Numbers](#)

There are other primitive integer types: `UInt8`, `UInt16`, `UInt32`, and `UInt64` are unsigned integers, e.g., we do not interpret the number as negative if the first bit is `1`. As they tend to be used to represent bit sequences they are displayed in hexadecimal, that is base-16, using digits 0-9a-c, e.g., $12 = (c)_{16}$:

```
In [8]: UInt16(12)
```

```
Out[8]: 0x000c
```

A non-primitive type is `BigInt` which allows arbitrary length integers.

2. Strings and parsing

We have seen that `bitstring` returns a string of bits. Strings can be created with quotation marks

```
In [9]: str = "hello world 😊"
```

```
Out[9]: "hello world 😊"
```

We can access characters of a string with brackets:

```
In [10]: str[1], str[13]
```

```
Out[10]: ('h', '😊')
```

Each character is a primitive type, in this case using 32 bits/4 bytes:

```
In [11]: typeof(str[6]), length(bitstring(str[6]))
```

```
Out[11]: (Char, 32)
```

Strings are not primitive types, but rather point to the start of a sequence of `Char`s in memory. In this case, there are $32 * 13 = 416$ bits/52 bytes in memory.

Strings are *immutable*: once created they cannot be changed. But a new string can be created that modifies an existing string. The simplest example is `*`, which concatenates two strings:

```
In [12]: "hi" * "bye"
```

```
Out[12]: "hibye"
```

(Why `*`? Because concatenation is non-commutative.) We can combine this with indexing to, for example, create a new string with a different last character:

```
In [13]: str[1:end-1] * "😂"
```

```
Out[13]: "hello world 😂"
```

Parsing strings

We can use the command `parse` to turn a string into an integer:

```
In [14]: parse{Int, "123"}
```

```
Out[14]: 123
```

We can specify base 2 as an optional argument:

```
In [15]: parse{Int, "-101"; base=2}
```

```
Out[15]: -5
```

If we are specifying bits its safer to parse as an `UInt32`, otherwise the first bit is not recognised as a sign:

```
In [16]: bits = "11110000100111111001100110001010"
x = parse{UInt32, bits; base=2}
```

```
Out[16]: 0xf09f998a
```

The function `reinterpret` allows us to reinterpret the resulting sequence of 32 bits as a different type. For example, we can reinterpret as an `Int32` in which case the first bit is taken to be the sign bit and we get a negative number:

```
In [17]: reinterpret{Int32, x}
```



```

@ Base ./array.jl:861
[2] top-level scope
@ In[22]:1
[3] eval
@ ./boot.jl:373 [inlined]
[4] include_string(mapexpr::typeof(REPL.softscope), mod::Module, code::String, file
name::String)
@ Base ./loading.jl:1196

```

Vectors can be made with different types, for example, here is a vector of three 8-bit integers:

```
In [23]: v = [Int8(11), Int8(24), Int8(32)]
```

```
Out[23]: 3-element Vector{Int8}:
 11
 24
 32
```

Just like strings, Vectors are not primitive types, but rather point to the start of sequence of bits in memory that are interpreted in the corresponding type. In this last case, there are $3 * 8 = 24$ bits/3 bytes in memory.

The easiest way to create a vector is to use `zeros` to create a zero `Vector` and then modify its entries:

```
In [24]: v = zeros{Int, 5}
v[2] = 3
v
```

```
Out[24]: 5-element Vector{Int64}:
 0
 3
 0
 0
 0
```

Note: we can't assign a non-integer floating point number to an integer vector:

```
In [25]: v[2] = 3.5
```

```
InexactError: Int64(3.5)
```

```
Stacktrace:
```

```

[1] Int64
@ ./float.jl:812 [inlined]
[2] convert
@ ./number.jl:7 [inlined]
[3] setindex!(A::Vector{Int64}, x::Float64, i1::Int64)
@ Base ./array.jl:903
[4] top-level scope
@ In[25]:1
[5] eval
@ ./boot.jl:373 [inlined]
[6] include_string(mapexpr::typeof(REPL.softscope), mod::Module, code::String, file
name::String)
@ Base ./loading.jl:1196

```

We can also create vectors with `ones` (a vector of all ones), `rand` (a vector of random numbers between 0 and 1) and `randn` (a vector of samples of normal distributed quasi-random numbers).

When we create a vector whose entries are of different types, they are mapped to a type that can represent every entry. For example, here we input a list of one `Int32` followed by three `Int64` s, which are automatically converted to all be `Int64` :

```
In [26]: [Int32(1), 2, 3, 4]
```

```
Out[26]: 4-element Vector{Int64}:
 1
 2
 3
 4
```

In the event that the types cannot automatically be converted, it defaults to an `Any` vector, which is similar to a Python list. This is bad performance-wise as it does not know how many bits each element will need, so should be avoided.

```
In [27]: [1.0, 1, "1"]
```

```
Out[27]: 3-element Vector{Any}:
 1.0
 1
 "1"
```

We can also specify the type of the Vector explicitly by writing the desired type before the first bracket:

```
In [28]: Int32[1, 2, 3]
```

```
Out[28]: 3-element Vector{Int32}:
 1
 2
 3
```

We can also create an array using comprehensions:

```
In [29]: [k^2 for k = 1:5]
```

```
Out[29]: 5-element Vector{Int64}:
 1
 4
 9
16
25
```

Matrices are created similar to vectors, but by specifying two dimensions instead of one. Again, the simplest way is to use `zeros` to create a matrix of all zeros:

```
In [30]: zeros{Int, 5, 5} # creates a 5x5 matrix of Int zeros
```

```
Out[30]: 5x5 Matrix{Int64}:
 0 0 0 0 0
 0 0 0 0 0
 0 0 0 0 0
 0 0 0 0 0
 0 0 0 0 0
```

We can also create matrices by hand. Here, spaces delimit the columns and semicolons delimit the rows:

```
In [31]: A = [1 2; 3 4; 5 6]
```

```
Out[31]: 3×2 Matrix{Int64}:
 1  2
 3  4
 5  6
```

We can also create matrices using brackets, a formula, and a `for` command:

```
In [32]: [k^2+j for k=1:5, j=1:5]
```

```
Out[32]: 5×5 Matrix{Int64}:
 2  3  4  5  6
 5  6  7  8  9
10 11 12 13 14
17 18 19 20 21
26 27 28 29 30
```

Matrices are really vectors in disguise. They are still stored in memory in a consecutive sequence of bits. We can see the underlying vector using the `vec` command:

```
In [33]: vec(A)
```

```
Out[33]: 6-element Vector{Int64}:
 1
 3
 5
 2
 4
 6
```

The only difference between matrices and vectors from the computers perspective is that they have a `size` which changes the interpretation of whats stored in memory:

```
In [34]: size(A)
```

```
Out[34]: (3, 2)
```

Matrices can be manipulated easily on a computer. We can multiply a matrix times vector:

```
In [35]: x = [8; 9]
         A * x
```

```
Out[35]: 3-element Vector{Int64}:
 26
 60
 94
```

or a matrix times matrix:

```
In [36]: A * [4 5; 6 7]
```

```
Out[36]: 3×2 Matrix{Int64}:
16 19
36 43
56 67
```

If you use `.*`, it does entrywise multiplication:

```
In [37]: [1 2; 3 4] .* [4 5; 6 7]
```

```
Out[37]: 2×2 Matrix{Int64}:
          4  10
          18  28
```

We can take the transpose of a real vector as follows:

```
In [38]: a = [1, 2, 3]
          a'
```

```
Out[38]: 1×3 adjoint(::Vector{Int64}) with eltype Int64:
          1  2  3
```

Note for complex-valued vectors this is the conjugate-transpose, and so one may need to use `transpose(a)`. Both `a'` and `transpose(a)` should be thought of as "dual-vectors", and so multiplication with a transposed vector with a normal vector gives a constant:

```
In [39]: b = [4, 5, 6]
          a' * b
```

```
Out[39]: 32
```

One important note: a vector is not the same as an $n \times 1$ matrix, and a transposed vector is not the same as a $1 \times n$ matrix.

Accessing and altering subsections of arrays

We will use the following notation to get at the columns and rows of matrices:

```
A[a:b,k]    # returns the a-th through b-th rows of the k-th column of
              A as a Vector of length (b-a+1)
A[k,a:b]    # returns the a-th through b-th columns of the k-th row of
              A as a Vector of length (b-a+1)
A[:,k]      # returns all rows of the k-th column of A as a Vector of
              length size(A,1)
A[k,:]      # returns all columns of the k-th row of A as a Vector of
              length size(A,2)
A[a:b,c:d]  # returns the a-th through b-th rows and c-th through d-th
              columns of A
              # as a (b-a+1) x (d-c+1) Matrix
```

The ranges `a:b` and `c:d` can be replaced by any `AbstractVector{Int}`. For example:

```
In [40]: A = [1 2 3; 4 5 6; 7 8 9; 10 11 12]
           A[[1,3,4],2] # returns the 1st, 3rd and 4th rows of the 2nd column of A
```

```
Out[40]: 3-element Vector{Int64}:
          2
          8
         11
```

Exercise Can you guess what `A[2,[1,3,4]]` returns, using the definition of `A` as above? What about `A[1:2,[1,3]]`? And `A[1,B[1:2,1]]`? And `vec(A[1,B[1:2,1]])`?

We can also use this notation to modify entries of the matrix. For example, we can set the $1:2 \times 2:3$ subblock of `A` to `[1 2; 3 4]` as follows:


```
In [41]: A[1:2,2:3] = [1 2; 3 4]
A
```

```
Out[41]: 4×3 Matrix{Int64}:
 1  1  2
 4  3  4
 7  8  9
10 11 12
```

Broadcasting

It often is necessary to apply a function to every entry of a vector. By adding `.` to the end of a function we "broadcast" the function over a vector:

```
In [42]: x = [1,2,3]
cos.(x) # equivalent to [cos(1), cos(2), cos(3)]
```

```
Out[42]: 3-element Vector{Float64}:
 0.5403023058681398
-0.4161468365471424
-0.9899924966004454
```

Broadcasting has some interesting behaviour for matrices. If one dimension of a matrix (or vector) is 1, it automatically repeats the matrix (or vector) to match the size of another example.

Example

```
In [43]: [1,2,3] .* [4,5]'
```

```
Out[43]: 3×2 Matrix{Int64}:
 4  5
 8 10
12 15
```

Since `size([1,2,3],2) == 1` it repeats the same vector to match the size `size([4,5]',2) == 2`. Similarly, `[4,5]'` is repeated 3 times. So the above is equivalent to:

```
In [44]: [1 1; 2 2; 3 3] .* [4 5; 4 5; 4 5]
```

```
Out[44]: 3×2 Matrix{Int64}:
 4  5
 8 10
12 15
```

Note we can also use broadcasting with our own functions (construction discussed later):

```
In [45]: f = (x,y) -> cos(x + 2y)
f.([1,2,3], [4,5]')
```

```
Out[45]: 3×2 Matrix{Float64}:
-0.91113  0.0044257
-0.839072 0.843854
 0.0044257 0.907447
```

Ranges

We have already seen that we can represent a range of integers via `a:b`. Note we can convert it to a `Vector` as follows:

```
In [46]: Vector{2:6}
```

```
Out[46]: 5-element Vector{Int64}:
 2
 3
 4
 5
 6
```

We can also specify a step:

```
In [47]: Vector{2:2:6}, Vector{6:-1:2}
```

```
Out[47]: ([2, 4, 6], [6, 5, 4, 3, 2])
```

Finally, the `range` function gives more functionality, for example, we can create 4 evenly spaced points between `-1` and `1`:

```
In [48]: Vector(range(-1, 1; length=4))
```

```
Out[48]: 4-element Vector{Float64}:
 -1.0
 -0.3333333333333333
  0.3333333333333333
  1.0
```

Note that `Vector` is mutable but a range is not:

```
In [49]: r = 2:6
r[2] = 3  # Not allowed
```

setindex! not defined for UnitRange{Int64}

Stacktrace:

```
[1] error(::String, ::Type)
  @ Base ./error.jl:42
[2] error_if_canonical_setindex(#unused#::IndexLinear, A::UnitRange{Int64}, #unused
#::Int64)
  @ Base ./abstractarray.jl:1323
[3] setindex!(A::UnitRange{Int64}, v::Int64, I::Int64)
  @ Base ./abstractarray.jl:1314
[4] top-level scope
  @ In[49]:2
[5] eval
  @ ./boot.jl:373 [inlined]
[6] include_string(mapexpr::typeof(REPL.softscope), mod::Module, code::String, file
name::String)
  @ Base ./loading.jl:1196
```

4. Types

Julia has two different kinds of types: primitive types (like `Int64`, `Int32`, `UInt32` and `Char`) and composite types. Here is an example of an in-built composite type representing complex numbers, for example, $z = 1 + 2i$:

```
In [50]: z = 1 + 2im
typeof(z)
```

Out[50]: Complex{Int64}

A complex number consists of two fields: a real part (denoted `re`) and an imaginary part (denoted `im`). Fields of a type can be accessed using the `.` notation:

In [51]: `z.re, z.im`

Out[51]: (1, 2)

We can make our own types. Let's make a type to represent complex numbers in the format

$$z = r \exp(i\theta)$$

That is, we want to create a type with two fields: `r` and `θ`. This is done using the `struct` syntax, followed by a list of names for the fields, and finally the keyword `end`.

In [52]:

```
struct RadialComplex
    r
    θ
end
z = RadialComplex(1,0.1)
```

Out[52]: RadialComplex(1, 0.1)

We can access the fields using `.`:

In [53]: `z.r, z.θ`

Out[53]: (1, 0.1)

Note that the fields are immutable: we can create a new `RadialComplex` but we cannot modify an existing one. To make a mutable type we use the command `mutable struct`:

In [54]:

```
mutable struct MutableRadialComplex
    r
    θ
end

z = MutableRadialComplex(1,2)
z.r = 2
z.θ = 3
z
```

Out[54]: MutableRadialComplex(2, 3)

Abstract types

Every type is a sub-type of an *abstract type*, which can never be instantiated on its own. For example, every integer and floating point number is a real number.

Therefore, there is an abstract type `Real`, which encapsulates many other types, including `Float64`, `Float32`, `Int64` and `Int32`.

We can test if type `T` is part of an abstract type `V` using the syntax `T <: V`:

```
In [55]: Float64 <: Real, Float32 <: Real, Int64 <: Real
```

```
Out[55]: (true, true, true)
```

Every type has one and only one super type, which is *always* an abstract type.

The function `supertype` applied to a type returns its super type:

```
In [56]: supertype(Int32) # returns Signed, which represents all signed integers.
```

```
Out[56]: Signed
```

```
In [57]: supertype(Float32) # returns `AbstractFloat`, which is a subtype of `Real`
```

```
Out[57]: AbstractFloat
```

An abstract type also has a super type:

```
In [58]: supertype(Real)
```

```
Out[58]: Number
```

Type annotation and templating

The types `RadialComplex` and `MutableRadialComplex` won't be efficient as we have not told the compiler the type of `r` and `θ`. For the purposes of this module, this is fine as we are not focussing on high performance computing. However, it may be of interest how to rectify this.

We can impose a type on the field name with `::`:

```
In [59]: struct FastRadialComplex
           r::Float64
           θ::Float64
       end
       z = FastRadialComplex(1,0.1)
       z.r, z.θ
```

```
Out[59]: (1.0, 0.1)
```

In this case `z` is stored using precisely 128-bits.

Sometimes we want to support multiple types. For example, we may wish to support 32-bit floats. This can be done as follows:

```
In [60]: struct TemplatedRadialComplex{T}
           r::T
           θ::T
       end
       z = TemplatedRadialComplex(1f0,0.1f0) # f0 creates a `Float32`
```

```
Out[60]: TemplatedRadialComplex{Float32}(1.0f0, 0.1f0)
```

This is stored in precisely 64-bits.

Relationship with C structs, heap and stack (advanced)

For those familiar with C, a `struct` in Julia whose fields are primitive types or composite types built from primitive types, is exactly equivalent to a `struct` in C, and can in fact be passed to C functions without any performance cost. Behind the scenes Julia uses the LLVM compiler and so C and Julia can be freely mixed.

Another thing to note is that there are two types of memory: the [stack](#) and the [heap](#). The stack has fixed memory length and is much faster as it avoids dynamic allocation and deallocation of memory. So an instance of a type with a known fixed length (like `FastRadialComplex`) will typically be in the stack and be much faster than an instance of a type with unknown or variable length (like `RadialComplex` or `Vector`), which will be on the heap.

For stack-allocated instances, the compiler may even go a step further and compile a function so that an instance of a type only lives on the cache or even in registers.

5. Loops and branches

For loops work essentially the same as in Python. The one caveat is to remember we are using 1-based indexing, e.g., `1:5` is a range consisting of `[1,2,3,4,5]` :

In [61]:

```
for k = 1:5
    println(k^2)
end
```

```
1
4
9
16
25
```

While loops look like

In [62]:

```
x = 1
while x < 5
    println("x is $x which is less than 5, incrementing!")
    x += 1
end
x
```

```
x is 1 which is less than 5, incrementing!
x is 2 which is less than 5, incrementing!
x is 3 which is less than 5, incrementing!
x is 4 which is less than 5, incrementing!
5
```

Out[62]:

If-elseif-else statements look like:

In [63]:

```
x = 5
if isodd(x)
    println("it's odd")
elseif x == 2
    println("it's 2")
else
```

```
println("it's even")
end
```

it's odd

6. Functions

Functions are created in a number of ways. The most standard way is using the keyword `function`, followed by a name for the function, and in parentheses a list of arguments.

Let's make a function that takes in a single number x and returns x^2 .

```
In [64]: function sq(x)
          x^2
        end
sq(2), sq(3)
```

Out[64]: (4, 9)

There is also a convenient syntax for defining functions on one line, e.g., we can also write

```
In [65]: sq(x) = x^2
```

Out[65]: sq (generic function with 1 method)

Multiple arguments to the function can be included with `,`.

Here's a function that takes in 3 arguments and returns the average.

(We write it on 3 lines only to show that functions can take multiple lines.)

```
In [66]: function av(x, y, z)
          ret = x + y
          ret = ret + z
          ret/3
        end
av(1, 2, 3)
```

Out[66]: 2.0

Variables live in different scopes. In the previous example, `x`, `y`, `z` and `ret` are *local variables*: they only exist inside of `av`.

So this means `x` and `z` are *not* the same as our complex number `x` and `z` defined above.

Warning: if you reference variables not defined inside the function, they will use the outer scope definition.

The following example shows that if we mistype the first argument as `xx`, then it takes on the outer scope definition `x`, which is a complex number:

```
In [67]: function av2(xx, y, z)
          (x + y + z)/3
        end
```

Out[67]: av2 (generic function with 1 method)

You should almost never use this feature!!

We should ideally be able to predict the output of a function from knowing just the inputs.

Example Let's create a function that calculates the average of the entries of a vector.

```
In [68]: function vecaverage(v)
           ret=0
           for k = 1:length(v)
               ret = ret + v[k]
           end
           ret/length(v)
       end
       vecaverage([1,5,2,3,8,2])
```

Out[68]: 3.5

Julia has an inbuilt `sum` command that we can use to check our code:

```
In [69]: sum([1,5,2,3,8,2])/6
```

Out[69]: 3.5

Functions with type signatures

functions can be defined only for specific types using `::` after the variable name. The same function name can be used with different type signatures.

The following defines a function `mydot` that calculates the dot product, with a definition changing depending on whether it is an `Integer` or a `Vector`.

Note that `Integer` is an abstract type that includes all integer types: `mydot` is defined for pairs of `Int64` 's, `Int32` 's, etc.

```
In [70]: function mydot(a::Integer, b::Integer)
           a*b
       end

       function mydot(a::AbstractVector, b::AbstractVector)
           # we assume length(a) == length(b)
           ret = 0
           for k = 1:length(a)
               ret = ret + a[k]*b[k]
           end
           ret
       end

       mydot(5, 6) # calls the first definition
```

Out[70]: 30

```
In [71]: mydot(Int8(5), Int8(6)) # also calls the first definition
```

Out[71]: 30

```
In [72]: mydot(1:3, [4,5,6]) # calls the second definition
```

Out[72]: 32

We should actually check that the lengths of `a` and `b` match.

Let's rewrite `mydot` using an `if`, `else` statement. The following code only does the for loop if the length of `a` is equal to the length of `b`, otherwise, it throws an error.

If we name something with the exact same signature (name, and argument types), previous definitions get overridden. Here we correct the implementation of `mydot` to throw an error if the lengths of the inputs do not match:

```
In [73]: function mydot(a::AbstractVector, b::AbstractVector)
           ret=0
           if length(a) == length(b)
               for k = 1:length(a)
                   ret = ret + a[k]*b[k]
               end
           else
               error("arguments have different lengths")
           end
           ret
       end
mydot([1,2,3], [5,6,7,8])
```

arguments have different lengths

Stacktrace:

```
[1] error(s::String)
   @ Base ./error.jl:33
[2] mydot(a::Vector{Int64}, b::Vector{Int64})
   @ Main ./In[73]:8
[3] top-level scope
   @ In[73]:12
[4] eval
   @ ./boot.jl:373 [inlined]
[5] include_string(mapexpr::typeof(REPL.softscope), mod::Module, code::String, file
name::String)
   @ Base ./loading.jl:1196
```

Anonymous (lambda) functions

Just like Python it is possible to make anonymous functions, with two variants on syntax:

```
In [74]: f = x -> x^2
          g = function(x)
              x^2
          end
```

Out[74]: #9 (generic function with 1 method)

There is not much difference between named and anonymous functions, both are compiled in the same manner. One can essentially think of named functions as "global constant anonymous functions".

Tuples

Tuples are similar to vectors but written with the notation `(x,y,z)` instead of `[x,y,z]`.

They allow the storage of *different types*. For example:


```
In [75]: t = (1,2.0,"hi")
```

```
Out[75]: (1, 2.0, "hi")
```

On the surface, this is very similar to a `Vector{Any}` :

```
In [76]: v=[1,2.0,"hi"]
```

```
Out[76]: 3-element Vector{Any}:
 1
 2.0
 "hi"
```

The main difference is that a `Tuple` knows the type of its arguments:

```
In [77]: typeof(t)
```

```
Out[77]: Tuple{Int64, Float64, String}
```

The main benefit of tuples for us is that they provide a convenient way to return multiple arguments from a function. For example, the following returns both `cos(x)` and `x^2` from a single function:

```
In [78]: function mytuplereturn(x)
           (cos(x), x^2)
       end
       mytuplereturn(5)
```

```
Out[78]: (0.28366218546322625, 25)
```

We can also employ the convenient syntax to create two variables at once:

```
In [79]: x,y = mytuplereturn(5)
```

```
Out[79]: (0.28366218546322625, 25)
```

Modules, Packages, and Plotting

Julia, like Python, has modules and packages. For example to load support for linear algebra functionality like `norm` and `det`, we need to load the `LinearAlgebra` module:

```
In [80]: using LinearAlgebra
          norm([1,2,3]), det([1 2; 3 4])
```

```
Out[80]: (3.7416573867739413, -2.0)
```

It is fairly straightword to create ones own modules and packages, however, we will not need modules in this....module.

Plotting


Some important functionality such as plotting requires non-built in packages. There are many packages such as [PyPlot.jl](#), which wraps Python's [matplotlib](#) and [Makie.jl](#), which is a state-of-the-art GPU based 3D plotting package. We will use [Plots.jl](#), which is an umbrella package that supports different backends.

For example, we can plot a simple function as follows:

In [81]:

```
using Plots
x = range(-1, 1; length=1000) # Create a range of a 1000 evenly spaced numbers between -1 and 1
y = cos.(x) # Create a new vector with `cos` applied to each entry of `x`
plot(x, y; label="cos(x)")
plot!(x, 2y; label="2cos(x)")
```

Out[81]:



Note the `!` is just a convention: any function that modifies its input or global state should have `!` at the end of its name.

Installing packages (advanced)

If you choose to use Julia on your own machine, you may need to install packages. This can be done by typing the following, either in Jupyter or in the REPL: `] add Plots`.